

Basic SQL

Instructor: Wang-Chiew Tan

Reference:
A First Course in Database Systems,
3rd edition, Chapter 2.3

9

What is a Data Model?

- A *data model* is a mathematical formalism that consists of two parts:
 1. A notation for describing data and mathematical objects for representing data.
 2. A set of operations for manipulating data.
- The relational data model
- What is the associated query language for the relational data model?

10

Two different query languages

- Codd proposed two different query languages for the relational data model.
 - Relational Algebra
 - Queries are expressed as a sequence of operations on relations.
 - Procedural language.
 - Relational Calculus
 - Queries are expressed as formulas of first-order logic.
 - Declarative language.
- **Codd's theorem:** The Relational Algebra query language has the same *expressive power* as the Relational Calculus query language.

11

Procedural vs. Declarative Languages

- Procedural program:
 - The program is specified as a sequence of operations to obtain the desired the outcome. I.e., *how* the outcome is to be obtained.
 - E.g., Java, C, ...
- Declarative program:
 - The program specifies *what* is the expected outcome, and not *how* the outcome is to be obtained.
 - E.g., Scheme, Ocaml, ...

12

An example

- In mathematical terms, the sequence F_n of Fibonacci numbers is defined by the recurrence relation

$$F_n = F_{n-1} + F_{n-2}$$

with seed values $F_0 = 0$ and $F_1 = 1$.

$$F(n) = \begin{cases} 0 & n = 0 \\ 1 & n = 1. \\ F(n-1) + F(n-2) & n > 1 \end{cases}$$

- 0,1,1,2,3,5,8,13,...

13

Fibonacci numbers in C++

```
#include <iostream>

// Fibonacci numbers, imperative style
int fibonacci(int iterations)
{
    int first = 0, second = 1; // seed values

    for (int i = 0; i < iterations - 1; ++i) {
        int sum = first + second;
        first = second;
        second = sum;
    }

    return first;
}

int main()
{
    std::cout << fibonacci(10) << "\n";
    return 0;
}
```

14

Fibonacci numbers in Ocaml

```
let rec fibonacci = function
| 0 -> 0
| 1 -> 1
| n when n > 1 -> fibonacci (n-1) + fibonacci (n-2)
```

- NOTE: it is possible to write a program in C++ in a “declarative” way similar to the above.

15

SQL – Structured Query Language

- Is SQL a procedural or a declarative language?
- Keep this question in mind as we learn SQL.
 - We will revisit this question later.
- SQL – pronounced as “sequel”
- Principle language used to describe and manipulate data stored in relational database systems.
- Neither relational algebra nor relational calculus proposed by Codd.
- Current standard: SQL-2011

16

SQL DDL and SQL DML

- Two main aspects to SQL:
 - Data Definition Language (DDL)
 - Sublanguage of SQL used to create, delete, modify the definition of tables and views.
 - For declaring database schemas.
 - Data Manipulation Language (DML)
 - Sublanguage of SQL that allows users to insert, delete, and modify rows of tables and pose queries.
 - For asking questions about the database and modifying the database.

17

Relations in SQL

- Three types of relations
 1. Stored relations (or tables)
 - These are tables that contain tuples and can be modified or queried.
 2. Views
 - Views are relations that are defined in terms of other relations but they are not stored. They are constructed only when needed.
 3. Temporary tables
 - These are (intermediate) tables that are constructed by the SQL execution engine during the processing of SQL queries and discarded when done.

18

Atomic Data Types

- CHAR(n): fixed-length string of up to n characters.
- VARCHAR(n): also a string of up to n characters.
- BIT(n)
- BIT VARYING(n)
- BOOLEAN: true/false/unknown.
- INT or INTEGER
 - Analogous to `int` in C.
- SHORTINT
 - Analogous to `short int` in C.
- FLOAT or REAL
- DOUBLE PRECISION
 - Analogous to `double` in C.
- DATE and TIME
 - Character strings of a special form.

19

Defining a table

```
CREATE TABLE Movies (  
    title          CHAR(100),  
    year          INT,  
    length        INT,  
    genre         CHAR(10),  
    studioName    CHAR(30),  
    producerC#    INT  
);
```

title	year	length	genre	studioName	producerC#
-------	------	--------	-------	------------	------------

20

Defining a table (cont'd)

```
CREATE TABLE MovieStar (  
  name          CHAR(30),  
  address       VARCHAR(255),  
  gender        CHAR(1),  
  birthdate     DATE  
);
```


name	address	gender	birthdate
------	---------	--------	-----------

21

Modifying Relation Schemas

- DROP TABLE Movies;
 - The entire table is removed from the database schema.
- ALTER TABLE MovieStar ADD phone CHAR(16);
 - Adds an attribute “phone” with type CHAR(16) to the table MovieStar.

name	address	gender	birthdate



name	address	gender	birthdate	phone

- ALTER TABLE MovieStar DROP birthdate;

name	address	gender	phone
------	---------	--------	-------

22

Default Values

```
CREATE TABLE MovieStar (  
    name          CHAR(30),  
    address       VARCHAR(255) DEFAULT "Hollywood",  
    gender        CHAR(1),  
    birthdate     DATE,  
);
```

- When a new tuple is inserted and no value is specified for the attribute address, then the value for this attribute will default to "Hollywood".
- If no default value is declared explicitly, then the value of the attribute will default to NULL.
 - NULL is a special value in SQL to represent unknown values.

23

More examples on default values

```
CREATE TABLE MovieStar (  
    name          CHAR(30),  
    address       VARCHAR(255) DEFAULT 'Hollywood',  
    gender        CHAR(1) DEFAULT '?',  
    birthdate     DATE DEFAULT '0000-00-00',  
);
```

```
ALTER TABLE MovieStar ADD phone CHAR(16) DEFAULT 'unlisted';
```

24

Keys



- A *key constraint* (or a *key* in short) of a relation schema R is a subset K of attributes of R such that
 1. For every instance r of R, every two distinct tuples of r must differ in their values from K.
 - Contrapositive: if two tuples agree on their values from K, then they must be the same tuple.
 2. Minimal: no proper subset of K has the above property.
- A *superkey* is a set of attributes of R that includes a key of R.
 - Fact. All keys are superkeys but some superkeys are not keys.

25

Examples

- Student(studentID, name, major, gender, avgGPA).
 - {studentID} is a key. It is also a superkey.
 - {studentID, name} is a superkey but not a key.
 - {studentID, name, major, gender, avgGPA} is a superkey but not a key.
- There can be multiple keys in general.
 - One key is chosen and define as the *primary key*, while the rest are *candidate keys*.
- Student(studentID, name, dob, major, gender, avgGPA)
 - {studentID}, {name, dob} are keys and also superkeys.
 - {studentID} is the primary key.
 - {name, dob} is the candidate key.
 - {name, dob, avgGPA} is a superkey.



26

Declaring keys

- Two ways to declare an attribute or set of attributes to be a key in the CREATE TABLE statement that defines a stored relation.
- If a single attribute is a key, it can be declared in two ways:

<pre>CREATE TABLE MovieStar (name CHAR(30) PRIMARY KEY, address VARCHAR(255), gender CHAR(1), birthdate DATE);</pre>	<pre>CREATE TABLE MovieStar (name CHAR(30), address VARCHAR(255), gender CHAR(1), birthdate DATE, PRIMARY KEY (name));</pre>
---	---

27

Declaring keys (cont'd)

- If the key consists of multiple attributes, then it can be declared as follows:

```
CREATE TABLE Movies (  
    title      CHAR(100),  
    year      INT,  
    length    INT,  
    genre     CHAR(10),  
    studioName CHAR(30),  
    producerC# INT,  
    PRIMARY KEY (title,year)  
);
```

28

PRIMARY KEY vs. UNIQUE

- In the previous examples, the keyword “PRIMARY KEY” can be replaced by “UNIQUE”.

```
CREATE TABLE MovieStar (  
    name      CHAR(30),  
    address   VARCHAR(255),  
    gender    CHAR(1),  
    birthdate DATE,  
    PRIMARY KEY (name)  
);
```

```
CREATE TABLE MovieStar (  
    name      CHAR(30),  
    address   VARCHAR(255),  
    gender    CHAR(1),  
    birthdate DATE,  
    UNIQUE (name)  
);
```

29

PRIMARY KEY vs. UNIQUE (cont'd)

```
CREATE TABLE MovieStar (  
    name      CHAR(30),  
    address   VARCHAR(255),  
    gender    CHAR(1),  
    birthdate DATE,  
    PRIMARY KEY (name)  
);
```

- None of the tuples in the MovieStar relation have null name values.
- Tuples are uniquely identified by their name values.
- There can be at most one primary key for a table.

```
CREATE TABLE MovieStar (  
    name      CHAR(30),  
    address   VARCHAR(255),  
    gender    CHAR(1),  
    birthdate DATE,  
    UNIQUE (name)  
);
```

- Tuples in MovieStar relation can contain null name values.
- Tuples in MovieStar relation with non-null name values are uniquely identified by their name values.
- There can be multiple unique constraints for a table in addition to a primary key.

30

PostgreSQL Meta Commands

- \l
 - List the databases.

After you connect to a database, you can perform the following:

- \d
 - List the tables of a database.
- \d table
 - List the columns of a table.

<http://www.postgresqlforbeginners.com/2010/11/interacting-with-postgresql.html>

31

SQL DDL and SQL DML

- Two main aspects to SQL:
 - Data Definition Language (DDL)
 - Sublanguage of SQL used to create, delete, modify the definition of tables and views.
 - For declaring database schemas.
 - Data Manipulation Language (DML)
 - Sublanguage of SQL that allows users to insert, delete, and modify rows of tables and pose queries.
 - For asking questions about the database and modifying the database.

Reference:
A First Course in Database Systems,
3rd edition, Chapter 6.1

32

Database schema for our running examples

- Let us assume we have the following database schema with five relation schemas.

Movies(title, year, length, genre, studioName, producerC#)

StarsIn(movieTitle, movieYear, starName)

MovieStar(name, address, gender, birthdate)

MovieExec(name, address, cert#, netWorth)

Studio(name, address, presC#)

33

A simple query in SQL

- Find all movies produced by Disney Studios in 1990.

SELECT *

FROM Movies

WHERE studioName = 'Disney' AND year = 1990;

34

Basic form of a SQL query

- Basic form:

```
SELECT [DISTINCT] c1, c2, ..., cm
FROM R1, R2, ..., Rn
[WHERE condition]
```

Diagram labels:

- Attribute names (points to c_1, c_2, \dots, c_m)
- Relation names (points to R_1, R_2, \dots, R_n)

- We will focus on one relation R_1 for now.
- What is the semantics (or meaning) of such a query?

35

A simple query in SQL

- Find all movies produced by Disney Studios in 1990.

```
SELECT *
FROM Movies
WHERE studioName = 'Disney' AND year = 1990;
```

Diagram label:

- The symbol "*" is a shorthand for all attributes of relations in the FROM clause. (points to the asterisk in the SELECT clause)

Movies

Title	Year	Length	Genre	studioName	producerC#
Pretty Woman	1990	119	true	Disney	999
Monster's Inc.	1990	121	true	Dreamworks	223
Jurassic Park	1998	145	NULL	Disney	675

36

A simple query in SQL (cont'd)

- Find all movies produced by Disney Studios in 1990.

```
SELECT *  
FROM Movies  
WHERE studioName = 'Disney' AND year = 1990;
```

Result

Title	Year	Length	Genre	studioName	producerC#
Pretty Woman	1990	119	true	Disney	999

37

An even simpler query in SQL

- Find all movies.

```
SELECT *  
FROM Movies
```

Equivalent to:

```
SELECT *  
FROM Movies  
WHERE true
```

Movies

Title	Year	Length	Genre	studioName	producerC#
Pretty Woman	1990	119	true	Disney	999
Monster's Inc.	1990	121	true	Dreamworks	223
Jurassic Park	1998	145	NULL	Disney	675

38

An even simpler query in SQL

- Find all movies.
`SELECT *`
`FROM Movies`

Result

Title	Year	Length	Genre	studioName	producerC#
Pretty Woman	1990	119	true	Disney	999
Monster's Inc.	1990	121	true	Dreamworks	223
Jurassic Park	1998	145	NULL	Disney	675

39

An even simpler query in SQL (cont'd)

- Return the title and year of all movies.

```
SELECT title, year  
FROM Movies
```

A Projection Query:

- Only a subset of attributes from the relation(s) in the FROM clause is selected.

Result

Title	Year
Pretty Woman	1990
Monster's Inc.	1990
Jurassic Park	1998

40

An even simpler query in SQL (cont'd)

- Return the title and year of all movies produced by Disney Studios in 1990.
SELECT title, year
FROM Movies
WHERE studioName = 'Disney' AND year = 1990;

Result

Title	Year
Pretty Woman	1990

41

An even simpler query in SQL

- Return the years of all movies with length less than 140.
SELECT year
FROM Movies
WHERE length < 140;

SELECT **DISTINCT** year
FROM Movies
WHERE length < 140;

Result

Year
1990
1990

Multiset or bag semantics

Result

Year
1990

Set semantics

42

Bags (or multisets) vs. Sets

- From your basic set theory
- Every element in a set is distinct
 - E.g., {2,4,6} is a set but {2,4,6,2,2} is not a set.
- A bag (or multiset) may contain repeated elements.
 - E.g., {{2,4,6}} is a bag. So is {{2,4,6,2,2}}.
 - Equivalently written as {{ 2[3],4[1],6[1]}}.
- The order among elements in a set or bag is not important
 - E.g., {2,4,6} = {4,2,6} = {6,4,2}
 - {{2,4,6,2,2}} = {{2,2,2,6,4}} = {{6,2,2,4,2}}.

43

Aliasing

- Allows you to rename the attributes of the result.
- Return the title and length of all movies as name and, resp., duration.
SELECT title AS name, length AS duration
FROM Movies

Result	name	duration
	Pretty Woman	119
	Monster's Inc.	121
	Jurassic Park	145

44

Aggregates in the SELECT clause

- Aggregates are allowed in the SELECT clause.
- Return the title and length of all movies as name and, resp., duration in hours.

```
SELECT title AS name, length * 0.016667 AS durationInHours  
FROM Movies
```

Result

name	durationInHours
Pretty Woman	1.98
Monster's Inc.	2.02
Jurassic Park	2.42

45

Constants in the result

- Constants can also added to the SELECT clause.
- Every tuple will have the same constant as specified in the SELECT clause.

```
SELECT title AS name, length * 0.016667 AS durationInHours, 'hrs.' as  
inHours  
FROM Movies
```

Result

name	durationInHours	inHours
Pretty Woman	1.98	hrs.
Monster's Inc.	2.02	hrs.
Jurassic Park	2.42	hrs.

46


More on the conditions in the WHERE clause

- WHERE studioName = 'Disney' AND year = 1990;
- Comparison operators:
 - =, <>, <, >, <=, >=
 - Equal, not equal, less than, greater than, less than or equal, greater than or equal
 - E.g., WHERE year <= 1990
- Logical connectives:
 - AND, OR, NOT
 - E.g., WHERE NOT (studioName = 'DISNEY' AND year <= 1990)
- Arithmetic expressions
 - +, -, *, /, etc.
 - E.g.,
WHERE ((length*0.01667) > 2 OR (length < 100)) AND year > 2000

47

More on the conditions in the WHERE clause (cont'd)

- In general, the WHERE clause consists of a boolean combination of conditions using logical connectives AND, OR, NOT.
- Each condition is of the form
expression op expression
 - *expression* is a column name, a constant, or an arithmetic or string expression
 - *op* is a comparison operator (=, <>, <, >, <=, >=, LIKE)

More on this later.

48

String comparisons

- Strings are compared in lexicographical order.
 - Let $a_1.a_2. \dots .a_n$ and $b_1.b_2. \dots b_m$ be two strings.
 - We have $a_1.a_2. \dots .a_n < b_1.b_2. \dots b_m$ if
 1. $a_1.a_2. \dots .a_n$ is a proper prefix of $b_1.b_2. \dots b_m$, or,
 2. For some $1 \leq i < n$, we have $a_1=b_1, a_2=b_2, \dots, a_{i-1}=b_{i-1}$ and $a_i < b_i$.
- Examples:
 - 'Pretty' < 'Pretty Woman',
 - 'butterfingers' < 'butterfly'

49

Pattern Matching in SQL

- LIKE operator
 - $s \text{ LIKE } p, s \text{ NOT LIKE } p$
 - s is a string, p is a pattern
- '%' (stands for 0 or more arbitrary chars)
- '_' (stands for exactly one arbitrary char)

```
SELECT *  
FROM Movies  
WHERE title LIKE 'Pretty%'
```

title LIKE "P_etty%"

50

Pattern Matching in SQL

- How do you match titles with an apostrophe?
 - E.g., 'Monster's Inc.'

```
SELECT *  
FROM Movies  
WHERE title LIKE 'Monster's Inc.'
```

Wrong!

```
WHERE title LIKE 'Monsters''s Inc.'  
WHERE title LIKE "" // matches "
```

51

Date and Time

- See Chapter 6.1.6.

52

SQL's three-valued logic

Students	studentID	name	major	gender	avgGPA
	112	Ann	Computer Science	F	NULL
	327	Bob	NULL	M	3.90
	835	Carl	Physics	M	4.00

- What is the result of the comparison “major = ‘Computer Science’ ” AND “avgGPA > 3” ?
- Three-valued logic TRUE, FALSE, UNKNOWN, where NULL is synonymous to UNKNOWN.
- If major is NULL, then “major=‘Computer Science’” evaluates to UNKNOWN.
- If avgGPA is NULL, then “avgGPA > 3” evaluates to UNKNOWN.

53

SQL's three-valued logic (cont'd)

p	q	$p \text{ OR } q$	$p \text{ AND } q$	$p = q$
True	True			
True	False			
True	Unknown			
False	True			
False	False			
False	Unknown			
Unknown	True			
Unknown	False			
Unknown	Unknown			

p	NOT p
True	
False	
Unknown	

54

SQL's three-valued logic (cont'd)

p	q	$p \text{ OR } q$	$p \text{ AND } q$	$p = q$
True	True	True	True	True
True	False	True	False	False
True	Unknown	True	Unknown	Unknown
False	True	True	False	False
False	False	False	False	True
False	Unknown	Unknown	False	Unknown
Unknown	True	True	Unknown	Unknown
Unknown	False	Unknown	False	Unknown
Unknown	Unknown	Unknown	Unknown	Unknown

p	NOT p
True	False
False	True
Unknown	Unknown

55

```
SELECT *
FROM Movies
WHERE major = 'Computer Science' AND avgGPA > 3.0;
```

- If the condition evaluates to UNKNOWN, then the tuple will not be returned.
- Both Ann and Bob will not be returned. In fact, none of the tuples will be returned.

56

Ordering the Output

```
SELECT [DISTINCT] <list of attributes>  
FROM   R1, R2, ..., Rn  
[WHERE condition]  
[ORDER BY <list of attributes>]
```

- ORDER BY presents the result in a sorted order.
- By default, the result will be ordered in ascending order. Otherwise, we state:
 - ORDER BY <list of attributes> DESC

57

ORDER BY

```
SELECT *  
FROM Movies  
WHERE studioName = 'Disney' AND year = 1990  
ORDER BY length, title;
```

- The result will list movies that satisfy the condition in the WHERE clause in ascending order of length, followed by title.
 - Shortest length movies will be listed first.
 - Among all movies with the same length, the movies will be sorted in ascending order of title.
- ORDER BY length, title DESC;
 - Longest length movies will be listed first.
 - Among all movies with the same length, the movies will be sorted in ascending order of title.

58

Meaning of an SQL query with one relation in the FROM clause

```
SELECT [DISTINCT] c1, c2, ..., cm  
FROM R1  
[WHERE condition]  
[ORDER BY <list of attributes>] [DESC]
```

- Let Result denote an empty collection.
- For every tuple t from R_1 ,
 - if t satisfies *condition* (i.e., condition evaluates to true), then add the tuple that consists of c_1, c_2, \dots, c_m components of t into Result.
- If DISTINCT is stated in the SELECT clause, remove duplicates in Result.
- If ORDER BY <list of attributes> exists, order the tuples in Result according to ORDER BY clause.
- Return Result.

59

Summary

- Declarative vs. Procedural
- Data Definition Language (DDL)
 - CREATE TABLE, DROP TABLE, ALTER TABLE, PRIMARY KEY, UNIQUE, DEFAULT
- Data Manipulation Language (DML)
 - Basic SQL
 - SELECT...FROM...WHERE, SELECT DISTINCT...FROM...WHERE
 - NULL values in SQL

60