

SQL

Instructor: Wang-Chiew Tan

Reference:

*A First Course in Database Systems,
3rd edition, Chapter 6.2*

Meaning of an SQL query with multiple relations in the FROM clause

```
SELECT [DISTINCT] c1, c2, ..., cm  
FROM R1, R2, ..., Rn ←  
[WHERE condition]  
[ORDER BY <list of attributes>] [DESC]
```

Suppose we now have more
than 1 relation in the FROM
clause.

- Let Result denote an empty collection.
- For every tuple t_1 from R_1 , t_2 from R_2 , ..., t_n from R_n
 - if t_1, \dots, t_n satisfy *condition* (i.e., condition evaluates to true), then add the resulting tuple that consists of c_1, c_2, \dots, c_m components of t into Result.
- If DISTINCT is stated in the SELECT clause, remove duplicates in Result.
- If ORDER BY <list of attributes> exists, order the tuples in Result according to ORDER BY clause.
- Return Result.

Database schema for our running examples

- Let us assume we have the following database schema with five relation schemas.

Movies(title, year, length, genre, studioName, producerC#)

StarsIn(movieTitle, movieYear, starName)

MovieStar(name, address, gender, birthdate)

MovieExec(name, address, cert#, netWorth)

Studio(name, address, presC#)

Products and Joins in SQL

```
SELECT *  
FROM Movies, StarsIn;
```

Movies	Title	Year	Length	Genre	studioName	producerC#
	Pretty Woman	1990	119	true	Disney	999
	Monster's Inc.	1990	121	true	Dreamworks	223
	Jurassic Park	1998	145	NULL	Disney	675
StarsIn	movieTitle	movieYear	starName			
	Pretty Woman	1990	Julia Roberts			
	Monster's Inc.	1990	John Goodman			

Products and Joins in SQL (cont'd)

```
SELECT *  
FROM Movies, StarsIn  
WHERE title = movietitle;
```

Movies	title	year	length	genre	studioName	producerC#
	Pretty Woman	1990	119	true	Disney	999
	Monster's Inc.	1990	121	true	Dreamworks	223
	Jurassic Park	1998	145	NULL	Disney	675
StarsIn	movieTitle	movieYear	starName			
	Pretty Woman	1990	Julia Roberts			
	Monster's Inc.	1990	John Goodman			

Disambiguating attributes

```
SELECT *  
FROM Movies, StarsIn  
WHERE title = 'Pretty Woman' AND year > 1995;
```

Movies	title	year	length	genre	studioName	producerC#
	Pretty Woman	1990	119	true	Disney	999
	Monster's Inc.	1990	121	true	Dreamworks	223
	Jurassic Park	1998	145	NULL	Disney	675
StarsIn	title	movieYear	starName			
	Pretty Woman	1990	Julia Roberts			
	Monster's Inc.	1990	John Goodman			

What if the first attribute is called "title"?

Disambiguating attributes (cont'd)

```
SELECT *  
FROM Movies, StarsIn  
WHERE StarsIn.title = 'Pretty Woman' AND year > 1995;
```

Movies	title	year	length	genre	studioName	producerC#
	Pretty Woman	1990	119	true	Disney	999
	Monster’s Inc.	1990	121	true	Dreamworks	223
	Jurassic Park	1998	145	NULL	Disney	675
StarsIn	title	movieYear	starName			
	Pretty Woman	1990	Julia Roberts			
	Monster’s Inc.	1990	John Goodman			

Tuple variables

```
SELECT *  
FROM Movies m, StarsIn s  
WHERE title = movietitle;
```

- m and s are tuple variables.
- m binds to a tuple in the Movies relation.
- s binds to a tuple in StarsIn relation.

Self Joins

```
SELECT *  
FROM StarsIn s1, StarsIn s2  
WHERE s1.movieYear = s2.movieYear;
```

StarsIn	title	movieYear	starName
	Pretty Woman	1990	Julia Roberts
	Monster's Inc.	1990	John Goodman

SQL Join Expressions

Reference:
A First Course in Database Systems,
3rd edition, Chapter 6.3.6 – 6.3.8

R(A,B,C) and S(C,D,E)

- R **CROSS JOIN** S;
 - Product of the two relations R and S.
 - Schema of resulting relation: (R.A,R.B,R.C,S.C,S.D,S.E).
 - Equivalent to:
SELECT *
FROM R, S;
- R **JOIN S ON** B=D **AND** A=E;
 - Selects only tuples from R and S where B=D and A=E.
 - Schema of the resulting relation: (R.A,R.B,R.C,S.C,S.D,S.E);
 - Equivalent to:
SELECT *
FROM R, S
WHERE B=D AND A=E;

SQL Join (cont'd)

R(A,B,C) and S(C,D,E)
SELECT R.A, R.B, R.C, S.D, S.E
FROM R JOIN S ON B=D AND A=E;

- Selects only tuples from R and S where B=D and A=E.
- Schema of the resulting relation: (A,B,C,D,E).

Natural Joins

R(A,B,C) and S(C,D,E)

- R **NATURAL JOIN** S;
 - Schema of the resulting relation: (A,B,C,D,E)
 - Equivalent to:
SELECT R.A, R.B, R.C, S.D, S.E
FROM R,S
WHERE R.C = S.C;

Outerjoins

R(A,B,C) and S(C,D,E)

- R **NATURAL FULL OUTER JOIN** S;
 - Schema of the resulting relation: (A,B,C,D,E).
- R **NATURAL LEFT OUTER JOIN** S;
 - Schema of the resulting relation: (A,B,C,D,E).
- R **NATURAL RIGHT OUTER JOIN** S;
 - Schema of the resulting relation: (A,B,C,D,E).
- R **FULL OUTER JOIN S** ON B = D;
 - Schema of the resulting relation: (A,B,C,D,E).
- R **FULL LEFT OUTER JOIN** ON B=D;
- R **FULL RIGHT OUTER JOIN** ON B=D;

Set and Bag Operations in SQL

Reference:
A First Course in Database Systems,
3rd edition, Chapter 6.2.5, 6.4.1, 6.4.2

- Set Union, Set Intersection, Set Difference
- Bag Union, Bag Intersection, Bag Difference
- Other set/bag operations
 - IN, op ANY, op ALL, EXISTS, NOT IN, NOT EXISTS
 - More on these later.

Set Union

R(A,B,C), S(A,B,C)

- Input to union must be *union-compatible*.
 - R and S have the same set of attributes and the corresponding attributes are of the same type.
- Output of union has the same schema as R or S.
- Meaning: Output consists of the set of all tuples of R and S.

(SELECT *	(SELECT *
FROM R)	FROM R
UNION	WHERE A > 10)
(SELECT *	UNION
FROM S);	(SELECT *
	FROM S
	WHERE B < 300);

Bag Union

R(A,B,C), S(A,B,C)

- Input to union must be *union-compatible*.
 - R and S have the same set of attributes and the corresponding attributes are of the same type.
- Output of union has the same schema as R or S.
- Meaning: Output consists of the collection of all tuples of R and S, including duplicate tuples.

(SELECT *	(SELECT *
FROM R)	FROM R
	WHERE A > 10)
UNION ALL	UNION ALL
(SELECT *	(SELECT *
FROM S);	FROM S
	WHERE B < 300);

Any difference between these two queries?

```
(SELECT DISTINCT *  
FROM R  
WHERE A > 10)  
UNION ALL  
(SELECT DISTINCT *  
FROM S  
WHERE B < 300);
```

```
(SELECT *  
FROM R  
WHERE A > 10)  
UNION  
(SELECT *  
FROM S  
WHERE B < 300);
```

Intersection, Difference

- Like union, intersection and difference are binary operators.
 - Input to intersection/difference operator consists of two relations R and S and they must be union-compatible.
 - Output has the same type as R or S.
- Set Intersection, Bag Intersection
 - `<Query1> INTERSECT <Query2>`, `<Query1> INTERSECT ALL <Query2>`
 - Meaning: Find all tuples that are common to both R and S.
- Set Difference, Bag Difference
 - `<Query1> EXCEPT <Query2>`, `<Query1> EXCEPT ALL <Query2>`
 - Find all tuples in R but not in S.
 - EXCEPT: takes the set of Q1 minus the set of Q2.
 - EXCEPT ALL: takes the bag of Q1 minus the bag of Q2.

Difference

- `<Query1> EXCEPT <Query2> EXCEPT <Query3>;` means
`(<Query1> EXCEPT <Query2>) EXCEPT <Query3>;`

Subqueries

Reference:
A First Course in Database Systems,
3rd edition, Chapter 6.3.

- A subquery is a query that is embedded in another query.
- Queries with UNION, INTERSECT, and EXCEPT have two subqueries.
- Subqueries can return used as part of a boolean expression.
 - A subquery returns a constant (or scalar value) which can be compared against another constant in the WHERE clause.
- Subqueries can return relations.
- Subqueries can appear in the FROM clause, followed by a tuple variable, which will represent a tuple in the result of the subquery.

Subqueries that return scalar values

Movies(title, year, length, genre, studioName, producerC#)

MovieExec(name, address, cert#, netWorth)

- Find all names of executives who produced the movie 'Star Wars'.

```
SELECT name
FROM Movies, MovieExec
WHERE title='Star Wars' AND producerC# = cert#;
```


Subqueries that return scalar values

Movies(title, year, length, genre, studioName, producerC#)

MovieExec(name, address, cert#, netWorth)

- Find all names of executives who produced the movie 'Star Wars'.

```
SELECT name
FROM Movies, MovieExec
WHERE title='Star Wars' AND producerC# = cert#;
```



Subqueries that return relations

```
SELECT name
FROM MovieExec
WHERE cert# IN (SELECT producerC#
                FROM Movies
                WHERE title = 'Star Wars');
```

- **IN, NOT IN**

Is this query equivalent to the one above?

```
SELECT name
FROM MovieExec, Movies
WHERE cert# = producerC# AND title = 'Star Wars';
```

Subqueries with subqueries

```
SELECT name
FROM MovieExec
WHERE cert# IN (SELECT producerC#
                FROM Movies
                WHERE (title, year) IN (SELECT movieTitle, movieYear
                                        FROM StarsIn
                                        WHERE starName = 'Harrison Ford'));

SELECT name
FROM MovieExec, Movies, StarsIn
WHERE cert# = producerC# AND title = movieTitle AND year = movieYear AND
      starName = 'Harrison Ford';
```

Correlated Subqueries

- In all the examples so far, the inner query is independent of the outer query.
- An inner query can also depend on the outer query.
- Find all titles of movies that have been used for two or more movies.

```
SELECT title
FROM Movies m
WHERE year < ANY (SELECT year
                  FROM Movies
                  WHERE title = m.title);
```

Correlation via
tuple variable m.

Checks that year is less than at least one of the answers returned by the subquery.
< ANY, <= ANY, > ANY, >= ANY, <> ANY, = ANY.
< ALL, <= ALL, > ALL, >= ALL, <> ALL, = ALL.

Correlated Subqueries (cont'd)

```
SELECT starName
FROM StarsIn s
WHERE EXISTS (SELECT *
              FROM StarsIn c
              WHERE s.movieTitle <> c.movieTitle AND
                    s.movieYear = c.movieYear AND
                    s.starName = c.starName);
```

Checks that the subquery returns a non-empty result.
EXISTS, NOT EXISTS.

Set Comparison Operators

- $x \text{ IN } Q$
 - Returns true if x occurs in the collection Q .
- $x \text{ NOT IN } Q$
 - Returns true if x does not occur in the collection Q .
- $\text{EXISTS } Q$
 - Returns true if Q is a non-empty collection.
- $\text{NOT EXISTS } Q$
 - Returns true if Q is an empty collection.
- $\text{NOT EXISTS } (Q1 \text{ EXCEPT } Q2)$
 - Returns true if $(Q1-Q2)$ is an empty collection.

Set Comparison Operators (cont'd)

- $op \text{ ANY}, op \text{ ALL}$, where op is one of $\{ <, <=, >, >=, <>, = \}$.
- $x \text{ op ANY } Q$
 - Returns true if there *exists* an element y in the result of Q such that $x \text{ op } y$ is true.
- $x \text{ op ALL } Q$
 - Returns true if *for every* element y in Q , we have that $x \text{ op } y$ is true.

Subqueries in the FROM clause

- Find the names of all movie executives who produce movies that Harrison Ford acted in.

```

SELECT name
FROM MovieExec m, (SELECT producerC#
                    FROM Movies, StarsIn
                    WHERE title = movieTitle AND
                          year = movieYear AND
                          starName = 'Harrison Ford') p
WHERE m.cert# = p. producerC#;

Movies(title, year, length, genre, studioName, producerC#)
StarsIn(movieTitle, movieYear, starName)
MovieExec(name, address, cert#, netWorth)
    
```

More examples

Customers

<u>sid</u>	cname	level	type	age
36	Cho	Beginner	snowboard	18
34	Luke	Inter	snowboard	25
87	Ice	Advanced	ski	20
39	Paul	Beginner	ski	33

Activities

<u>sid</u>	<u>slopeid</u>	day
36	s3	01/05/13
36	s1	01/06/13
36	s1	01/07/13
87	s2	01/07/13
87	s1	01/07/13
34	s2	01/05/13

Slopes

<u>slopeid</u>	name	color
s1	Mountain Run	blue
s2	Olympic Lady	black
s3	Magic Carpet	green
s4	KT-22	black

Example 1

- Find the names of customers who went on some slope on 01/07/13.

```
SELECT c.cname
FROM Customers c, Activities a
WHERE a.day='01/07/13' AND a.sid = c.sid;
```

```
SELECT cname
FROM Customers c
WHERE c.sid IN (SELECT a.sid
                FROM Activities a
                WHERE a.day='01/07/13');
```

Example 2

- Find the names of customers who did not go on any slope on 01/07/13.

```
SELECT cname
FROM Customers c
WHERE c.sid NOT IN (SELECT a.sid
                    FROM Activities a
                    WHERE a.day='01/07/13');
```


Example 3

- Find all names of customers who went on the slope “Olympic Lady” on 01/07/13.

```
SELECT c.cname
FROM Customers c
WHERE c.id IN (SELECT a.id
               FROM Activities a
               WHERE a.day='01/07/13' AND a.slopeid IN ( SELECT s.slopeid
                                                         FROM Slopes s
                                                         WHERE s.name='Olympic lady'));
```

Example 4

- Determine the colors of all slopes that Cho went.

```
SELECT s.color
FROM Activities a, Slopes s
WHERE a.slope-id = s.slope-id AND
      a.id = (SELECT id
              FROM Customers
              WHERE cname='Cho');
```

Example 5

- Find the names of customers who have been on some slope on the day 01/07/13.

```
SELECT c.cname
FROM Customers c
WHERE EXISTS (SELECT *
              FROM Activities a
              WHERE a.sid = c.sid AND a.day='01/07/13');
```

Example 6

- Find the names of all customers who went on some slope.

```
SELECT c.cname
FROM Customers c
WHERE c.sid = ANY (SELECT sid
                  FROM Activities a);
```

Example 7

- Find the names of all skiers whose age is greater than every snowboarder.

```
SELECT c.name
FROM Customers c
WHERE c.type='skier' AND c.age >ALL (SELECT c.age
                                     FROM Customers c
                                     WHERE c.type = 'snowboard');
```

What happens if this subquery returns an empty set?



- Find the names of the oldest customers.

```
SELECT c.cname
FROM Customer c
WHERE c.age >= ALL (SELECT c.age
                   FROM Customers c);
```

Practice homework 3

The example database records information on bars, customers, beers, and the associations among them.

- Beers(name, manf): stores information about beers, including the manufacturer of each beer.
- Bars(name, city, addr, license, phone): stores information about bars including their city, street address, phone number and their operating license.
- Drinkers(name, city, addr, phone): stores information about drinkers, including their city, street address and phone number.
- Likes (drinker, beer): indicates which drinker likes which beers (note that a drinker may like many beers and many drinkers may like the same beer).
- Sells (bar, beer, price): indicates the price of each beer sold at each bar (note that each bar can sell many beers and many bars can sell the same beer, at possibly different prices).
- Frequents (drinker, bar): indicates which drinker frequents which bars (note that each drinker may frequent many bars and many drinkers may frequent the same bar).

1. Find the names of all beers, and their prices, served by the bar 'Blue Angel'.
2. Find the name and phone number of every drinker who likes the beer 'Budweiser'.
3. Find all bars frequented by both 'Vince' and 'Herb'.
4. Find all bars in 'Chicago' (and display all attributes) for which we know either the address (i.e., addr in our schema) or the phone number but not both.

END